

THE HONG KONG UNIVERSITY OF SCIENCE & TECHNOLOGY

Department of Computer Science and Engineering

COMP4331: Introduction to Data Mining

Fall 2022 Assignment 1

Due time and date: 11:59pm, October 9 (Sun), 2022.

IMPORTANT NOTES

- **Your grade will be based on the correctness and clarity.**
- **Late submission: 25 marks will be deducted for every 24 hours after the deadline.**
- **ZERO-Tolerance on Plagiarism: All involved parties will get zero mark.**

You are given a student performance dataset, which can be downloaded from the UCI Machine Learning Repository:

<https://archive.ics.uci.edu/ml/machine-learning-databases/forest-fires/>

Please **DO** use the data “forestfires.csv” for the following tasks. A notebook template is provided for you on Google Colab:

https://colab.research.google.com/drive/1_-4zEbe08j-yhCMv5NiAWqUG2bTmjN5e?usp=sharing

Tasks:

1. Report the mean, standard deviation, mode, median, and the five-number summary of the attributes *RH* and *wind*.
2. For the attribute *RH*, use `seaborn` or `matplotlib` to show boxplots for
 - (a) *RH*;
 - (b) *RH* for various values of *month* (show *month* in x-axis);
 - (c) *RH* for various values of *day* (show *day* in x-axis).
3. For attributes *wind* and *RH*, use `seaborn` or `matplotlib` to show their histograms (use 8 equal-sized bins).
4. Use `seaborn` or `matplotlib` to show the scatter plot for attributes *wind* and *area* (show *area* in x-axis). Report the correlation coefficient between *wind* and *area*.
5. Consider the 6 attributes $\{FFMC, DMC, DC, ISI, RH, wind\}$, report the top 3 attributes that are most correlated (either positively or negatively) with *area*.
6. χ^2 -test (using the χ^2 table shown on the last page). Show your steps clearly (including contingency table, χ^2 -value, and *p*-value) in the report.
 - (a) By performing the χ^2 -test at a significance level of 0.01, are the attributes *temp* and *wind* independent of each other?
 - (b) By performing the χ^2 -test at a significance level of 0.01, are the attributes *X* and *Y* independent of each other?

7. Normalization:

- (a) Normalize attribute *FFMC* to the range $[0, 1]$ using min-max normalization. You can use the built-in function from `scikit-learn`.
- (b) Normalize attributes $\{DC, ISI, temp, RH, wind\}$ to mean zero and standard deviation one using z-score normalization. You can use built-in function from `scikit-learn`.
- (c) Output your results in parts (a) and (b) above to the csv file `data_normalized.csv` (use “,” as field delimiter and include column names in the header).

8. PCA: In this question, use only the attributes $\{DC, ISI, temp, RH, wind\}$ after the normalization in Task 7b.

- (a) Plot the cumulative explained variance ratio with the number of principal components.
- (b) Transform the data by PCA, by using the smallest number of PCA components such that the proportion of explained variance is at least 0.9. You can use the built-in PCA function from `scikit-learn`.
- (c) For the transformed data, name the new features *PC1*, *PC2*, and so on, and output the transformed dimensions to the csv file `data_reduced.csv` without the header (use “,” as field delimiter). Report the five-number summary for each obtained dimension.

9. Consider the original data. At first glance, this dataset has no missing value. However, for attribute *area*, some values got zero. We assume that these zeros can be viewed as missing values. Please fill in the missing values with the attribute mean of transactions whose corresponding attribute has non-zero value (filter non-zero value transactions first, then calculate attribute mean of them). and show the box plot for attribute *area* again. You can use the “`DataFrame.replace`” function from `pandas`. After filling in the missing values, report the correlation coefficient between *wind* and *area* again, and compare it with the value obtained in Task 4.

Submission Guidelines

Please submit (i) a report `report.pdf` for the answers for Q1 to Q6, Q8 ((a) and 5-number summary in (c)), and Q9 (the box plot and correlation coefficient analysis). Your steps in Q6 should be clear, (ii) a python notebook `assignment1.ipynb` for your code, and (iii) the output data files `data_normalized.csv` and `data_reduced.csv`. Zip all the files to `YourStudentID_assignment1.zip` (e.g., `12345678_assignment1.zip`). Please submit the assignment by uploading the compressed file to Canvas. Note that the assignment should be clearly legible, otherwise you may lose some points if the assignment is difficult to read. **Plagiarism will lead to zero point on this assignment.**

Degrees of Freedom	Probability										
	0.95	0.90	0.80	0.70	0.50	0.30	0.20	0.10	0.05	0.01	0.001
1	0.004	0.02	0.06	0.15	0.46	1.07	1.64	2.71	3.84	6.64	10.83
2	0.10	0.21	0.45	0.71	1.39	2.41	3.22	4.60	5.99	9.21	13.82
3	0.35	0.58	1.01	1.42	2.37	3.66	4.64	6.25	7.82	11.34	16.27
4	0.71	1.06	1.65	2.20	3.36	4.88	5.99	7.78	9.49	13.28	18.47
5	1.14	1.61	2.34	3.00	4.35	6.06	7.29	9.24	11.07	15.09	20.52
6	1.63	2.20	3.07	3.83	5.35	7.23	8.56	10.64	12.59	16.81	22.46
7	2.17	2.83	3.82	4.67	6.35	8.38	9.80	12.02	14.07	18.48	24.32
8	2.73	3.49	4.59	5.53	7.34	9.52	11.03	13.36	15.51	20.09	26.12
9	3.32	4.17	5.38	6.39	8.34	10.66	12.24	14.68	16.92	21.67	27.88
10	3.94	4.86	6.18	7.27	9.34	11.78	13.44	15.99	18.31	23.21	29.59
	Nonsignificant								Significant		